



# UPPMAX Introduction

2017-10-23

Martin Dahlö

[martin.dahlo@scilifelab.uu.se](mailto:martin.dahlo@scilifelab.uu.se)

Valentin Georgiev

[valentin.georgiev@icm.uu.se](mailto:valentin.georgiev@icm.uu.se)

Enabler for Life Sciences

**What is UPPMAX what it provides**

Jobs and queuing systems

Projects at UPPMAX

How to access UPPMAX

How to use the resources of UPPMAX

How to use the resources of UPPMAX in a good way!  
**Efficiency!!!**

## Uppsala Multidisciplinary Center for Advanced Computational Science

<http://www.uppmax.uu.se>

### 3 clusters

Rackham, 334 computer à 20 cores (128GB RAM)

Milou, 208 computers à 16 cores (128GB RAM)

17 with 256, 17 with 512

Bianca, 200 nodes à 16 cores (128GB RAM)

~11 PB fast parallel storage


Bioinformatics software

# The basic structure of supercomputer



Login nodes

## The basic structure of supercomputer




Calculation nodes




Login nodes

## The basic structure of supercomputer



Calculation nodes



Storage



Login nodes

**What is UPPMAX what it provides**

**Jobs and queuing systems**

Projects at UPPMAX

How to access UPPMAX

How to use the resources of UPPMAX

How to use the resources of UPPMAX in a good way!  
**Efficiency!!!**



# Job (computing)

From Wikipedia, the free encyclopedia

*For other uses, see [Job \(Unix\)](#) and [Job stream](#).*

In [computing](#), a **job** is a unit of work or unit of execution (that performs said work). A component of a job (as a unit of work) is called a *[task](#)* or a *[step](#)* (if sequential, as in a [job stream](#)). As a unit of execution, a job may be concretely identified with a single [process](#), which may in turn have subprocesses ([child processes](#); the process corresponding to the job being the [parent process](#)) which perform the tasks or steps that comprise the work of the job; or with a [process group](#); or with an abstract reference to a process or process group, as in [Unix job control](#).

Read/open files

Do something with the data

Print/save output


Read/open files

Do something with the data

Print/save output

## The basic structure of a supercomputer

Parallel computing is key  
Not one super fast



## The basic structure of a supercomputer

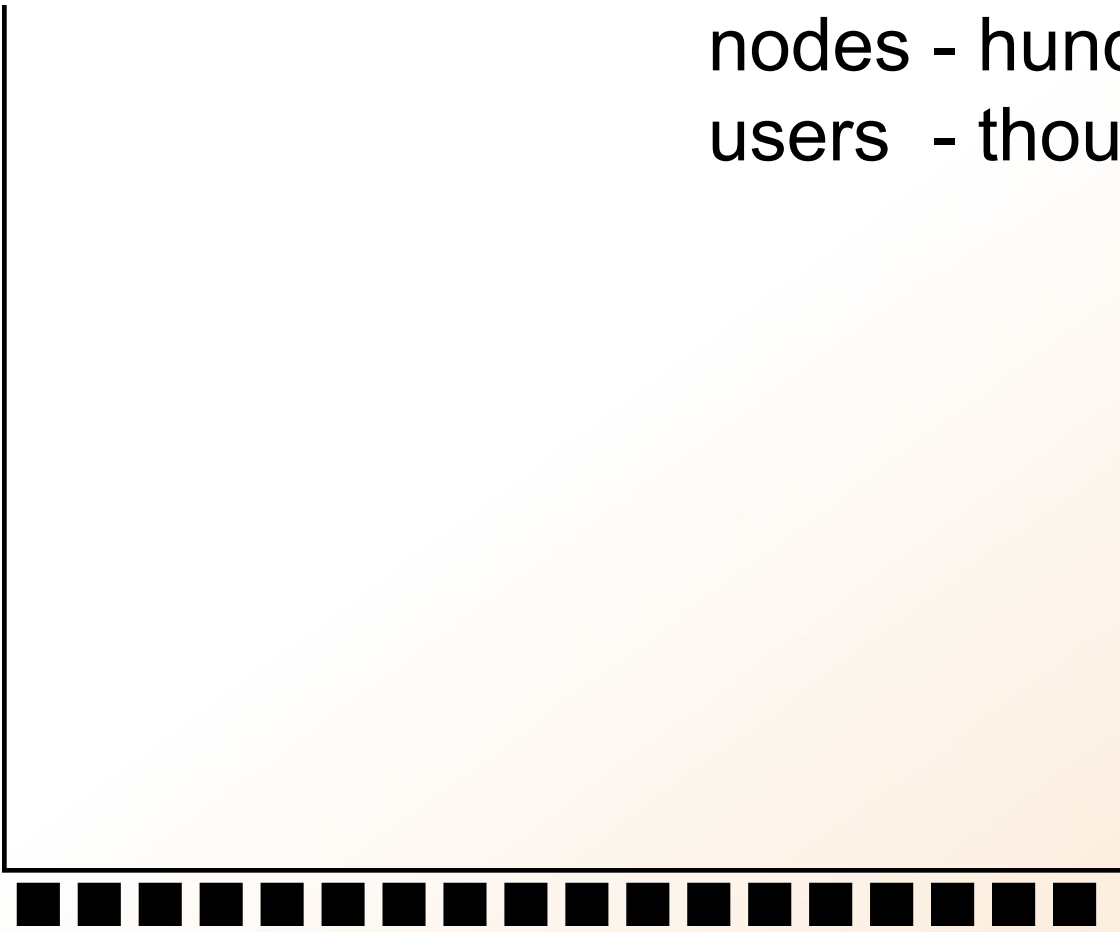
Parallel computing is key  
Not one super fast




# Queue System

More users than nodes  
Need for a queue


nodes - hundreds  
users - thousands



More users than nodes  
Need for a queue




More users than nodes  
Need for a queue





More users than nodes  
Need for a queue



**What is UPPMAX what it provides**

**Jobs and queuing systems**


**Projects at UPPMAX**

How to access UPPMAX

How to use the resources of UPPMAX

How to use the resources of UPPMAX in a good way!  
**Efficiency!!!**





Resources:

**compute**  
(core-hours/month)

**storage**  
(TB)

Today - UPPNEX projects:

cluster **Milou**

**2000** core-hours/month

**1 TB**

This course's project ID: g2017024

from 2018 - two separate projects:

SNIC project:

cluster **Rackham**

**2000** core-hours/month

**128 GB**

Scilifelab Storage project:

storage system **CREX**

**1 - 100 TB**

**What is UPPMAX what it provides**

**Jobs and queuing systems**

**Projects at UPPMAX**

**How to access UPPMAX**

How to use the resources of UPPMAX

How to use the resources of UPPMAX in a good way!  
**Efficiency!!!**



# How to access UPPMAX

SSH to a cluster

```
ssh -Y your_username@cluster_name.uppmax.uu.se
```

# SSH to Milou

```
VG-MBP:~ valentinggeorgiev$ ssh -Y valent@milou.uppmass.uu.se
Last login: Sun Oct 22 10:02:44 2017 from host-95-195-196-83.mobileonline.telia.com
```

[illegible]

#####

User Guides: <http://www.uppmax.uu.se/support/user-guides>

FAQ: <http://www.uppmax.uu.se/support/faq>

Write to [support@uppmax.uu.se](mailto:support@uppmax.uu.se), if you have questions or comments.

```
[valent@milou2 ~]$
```

# SSH to Rackham

```
VG-MBP:~ valentinggeorgiev$ ssh -Y valent@rackham.uppmass.uu.se
Last login: Sun Oct 22 10:14:21 2017 from host-95-195-196-83.mobileonline.telia.com
```

```
|_|_|_|_\\|_|_\\|_|_\\|_|_/\\|_|_\\_/_/|System:    rackham1
|_|_|_|_|_)|_|_)|_|_/_||_|/_|_|_|_\\|_|_/_/|User:      valent
|_|_|_|_|_/_|_|_/_|_|_|_|_|/_|_|_|_\\|_|_/_/|Jobs:     0 running
|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|_|Queue:    0 pending
```

#####

User Guides: <http://www.uppmax.uu.se/support/user-guides>

FAQ: <http://www.uppmax.uu.se/support/faq>

Write to [support@uppmax.uu.se](mailto:support@uppmax.uu.se), if you have questions or comments.


```
[valent@rackham1 ~]$
```




Local  
computer



Login  
nodes



Computation  
nodes



## Login nodes

use them to access UPPMAX

never use them to run **jobs**

don't even use them to do "quick stuff"

## Calculation nodes

do your work here - testing and running

## Calculation nodes

not accessible directly

SLURM (queueing system) gives you access

**What is UPPMAX what it provides**

**Jobs and queuing systems**

**Projects at UPPMAX**

**How to access UPPMAX**

**How to use the resources of UPPMAX**

How to use the resources of UPPMAX in a good way!  
**Efficiency!!!**



**workload manager**  
**job queue**  
**batch queue**  
**job scheduler**

**SLURM** (Simple Linux Utility for Resource Management)  
free and open source

- 1) Ask for node/core and run jobs manually**  
mainly for testing and small jobs
- 2) Write a script and submit it to SLURM**  
do the real job

## 1) Ask for node/core and run jobs manually

submit a request for resources



ssh to a calculation node



run programs

## 1) Ask for node/core and run jobs manually

```
salloc -A b2015245 -p core -n 1 -t 00:05:00
```

**salloc** - command

mandatory job parameters:

- A**        - project ID (who “pays”)
- p**        - node or core
- n**        - number of nodes/cores
- t**        - time

- A**      this course project g2017024  
            you have to be a member
  
- p**      1 node = 16 cores  
            1 hour walltime = 16 core-hours
  
- n**      number of cores (default value = 1)
- N**      number of nodes
  
- t**      format - hh:mm:ss  
            default value= 7-00:00:00  
            jobs killed when time limit reaches - always overestimate ~ 50%

## Information about your jobs

`squeue -u <user>`

```
[valent@milou2 valent]$ squeue -u valent
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
11334919	core	sh	valent	R	0:11	1	m164

```
[[valent@milou2 valent]$ salloc -A b2015245 -p core -n 1 -t 00:05:00 &
[2] 10994
```

```
[[valent@milou2 valent]$ salloc: Granted job allocation 11334919
```

```
[[valent@milou2 valent]$ squeue -u valent
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
11334919	core	sh	valent	R	0:11	1	m164

```
[[valent@milou2 valent]$ ssh -Y m164
```

_ \   _ \   \ /     / \ \ \ / /	System:	m164
)     )   \ /     / _ \ \ /	User:	valent
_     _ /   _ /         / _ _ \ / \	Jobs:	1 running
\ _ _ /   _     _     _ / _ \ \ / _ \	Queue:	0 pending

#####

User Guides: <http://www.uppmax.uu.se/support/user-guides>

FAQ: <http://www.uppmax.uu.se/support/faq>


Write to [support@uppmax.uu.se](mailto:support@uppmax.uu.se), if you have questions or comments.

```
[[valent@m164 ~]$ █
```

SSH to a calculation node (from a login node)

```
ssh -Y <node_name>
```





```
interactive -A b2015245 -p core -n 1 -t 00:05:00
```

```
System:      m1
User:        valent
Jobs:        1 running
Queue:       0 pending
```

#####

User Guides: <http://www.uppmax.uu.se/support/user-guides>  
FAQ: <http://www.uppmax.uu.se/support/faq>

Write to [support@uppmax.uu.se](mailto:support@uppmax.uu.se), if you have questions or comments.

```
[valent@m1 valent]$
```

## **2)Write a script and submit it to SLURM**

put all commands in a text file - script



tell SLURM to run the script  
(use the same job parameters)

## 2) Write a script and submit it to SLURM

put all commands in a text file - script

```
#!/bin/bash -l  
#SBATCH -A g2012157  
#SBATCH -p core  
#SBATCH -J Template_script  
#SBATCH -t 08:00:00
```

job parameters

```
# go to some directory  
cd ~/glob  
  
# do something  
echo Hello world!
```

tasks to be done

## 2) Write a script and submit it to SLURM

put all commands in a text file - script

```
#!/bin/bash -l
#SBATCH -A g2012157
#SBATCH -p node
#SBATCH -J Template_script
#SBATCH -t 08:00:00
```

```
# go to the correct directory
cd /home/dahlo/glob/work/uppmaxScripts/misc

# run tophat on the data, using 8 cores
tophat -p 8 /bubo/proj/g2012157/indexes/bowtie/hg19 tophat/input/ad12.fq
```

## **2)Write a script and submit it to SLURM**

tell SLURM to run the script  
(use the same job parameters)

**sbatch test.sbatch**

## 2) Write a script and submit it to SLURM

tell SLURM to run the script  
(use the same job parameters)

```
sbatch test.sbatch
```

**sbatch** - command

**test.sbatch** - name of the script file

## 2) Write a script and submit it to SLURM

tell SLURM to run the script  
(use the same job parameters)

```
sbatch -A b2015245 -p core -n 1 -t 00:05:00 test.sbatch
```



## Prints to a file instead of terminal slurm-<job id>.out

```
[valent@milou2 temp]$ ll
total 32
-rw-rw-r-- 1 valent valent 209 Oct 22 13:34 test.sbatch
[valent@milou2 temp]$ sbatch test.sbatch
Submitted batch job 11334939
[valent@milou2 temp]$ ll
total 64
-rw-rw-r-- 1 valent valent 31 Oct 22 13:35 slurm-11334939.out
-rw-rw-r-- 1 valent valent 209 Oct 22 13:34 test.sbatch
[valent@milou2 temp]$ cat slurm-11334939.out
this goes to slurm-<jobID>.out
[valent@milou2 temp]$ cat test.sbatch
#!/bin/bash -l

#SBATCH -A b2015245
#SBATCH -p core
#SBATCH -n 1
#SBATCH -t 00:05:00

# go to dir work
cd ~/work
# do something useless
echo "this goes to slurm-<jobID>.out"
echo "Hello, world!" > hello.txt
[valent@milou2 temp]$
```

Shows information about your jobs

`squeue -u <user>`

```
[valent@milou2 temp]$ sbatch test.sbatch
Submitted batch job 11334948
[valent@milou2 temp]$ squeue -u valent
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
11334948	core	test.sba	valent	CG	0:01	1	m200

`jobinfo -u <user>`

SLURM user guide

go to <http://www.uppmax.uu.se/>

click Support (left-hand side menu)

click User Guides

click Slurm user guide

or just google “uppmax slurm user guide”

link:

<http://www.uppmax.uu.se/support/user-guides/slurm-user-guide/>

100+ programs installed

Managed by a 'module system'

Installed, but hidden

Manually loaded before use

*module avail* - Lists all available modules

*module load <module name>* - Loads the module

*module unload <module name>* - Unloads the module

*module list* - Lists loaded modules

*module spider <word>* - Searches all modules after 'word'

Most bioinfo programs hidden under bioinfo-tools  
Load bioinfo-tools first, then program module

```
[dahlo@kalkyl3 work]$ module load cufflinks/1.2.1
ModuleCmd_Load.c(200):ERROR:105: Unable to locate a modulefile for 'cufflinks/1.2.1'
[dahlo@kalkyl3 work]$ module load bioinfo-tools
[dahlo@kalkyl3 work]$ module load cufflinks/1.2.1
[dahlo@kalkyl3 work]$
```

or

```
[dahlo@kalkyl3 work]$ module load samtools
ModuleCmd_Load.c(200):ERROR:105: Unable to locate a modulefile for 'samtools'
[dahlo@kalkyl3 work]$ module load bioinfo-tools samtools
[dahlo@kalkyl3 work]$
```

```
[dahlo@kalkyl4 work]$ module load bioinfo-tools
[dahlo@kalkyl4 work]$ module avail
```

```
----- /bubo/sw/mf/kalkyl/bioinfo-tools/alignment -----
MUMmer/3.22(default)      blast/2.2.24(default)      maq/0.7.1(default)
anfo/0.97                  blast/2.2.24+              mosaik-aligner/1.0.1388(default)
anfo/0.98(default)        blast/2.2.25              mosaik-aligner/1.1.0021
blast/2.2.15              blat/34                    mpiblast/1.6.0(default)
blast/2.2.18              bwa/0.5.8a                splitseek/1.3.2
blast/2.2.23              bwa/0.5.9                 splitseek/1.3.4(default)
blast/2.2.23+            hmmer/3.0
```

```
----- /bubo/sw/mf/kalkyl/bioinfo-tools/assembly -----
Ray/0.0.4                  abyss/1.2.4                abyss/1.3.0                velvet/1.0.03(default)
Ray/0.0.7(default)        abyss/1.2.5(default)        abyss/1.3.2                velvet/1.1.04
Ray/1.6.1                  abyss/1.2.7                mira/3.0.0                velvet/1.1.04_K101
abyss/1.2.3                abyss/1.2.7-maxk96          mira/3.2.0(default)        velvet/1.1.07
```

```
----- /bubo/sw/mf/kalkyl/bioinfo-tools/misc -----
BclConverter/1.7.1        freebayes/0.8.9            samtools/0.1.12-10(default)
BioPerl/1.6.1             freebayes/0.9.4            samtools/0.1.16
BioPerl/1.6.1_PERL5.10.1(default) gcta/0.92.0                samtools/0.1.18
BioPerl/1.6.1_PERL5.12.3  gcta/0.92.6                samtools/0.1.7a
FastQC/0.6.1              genomertools/1.3.5(default) samtools/0.1.8
FastQC/0.7.2(default)    htseq/0.4.6                samtools/0.1.9
Fastx/0.0.13(default)     htseq/0.5.1                snpEff/2.0.3
IGV/1.5.51                matrix2png/1.2.1            trinity/2011-05-13
biopython/1.56            picard/1.40                 trinity/2011-10-29
cellprofiler/20111024     picard/1.41
emmax/beta-07Mar2010      plink/1.07
```

```
----- /bubo/sw/mf/kalkyl/bioinfo-tools/phylogeny -----
concatpillar/1.4          garli/2.0                   raxml/7.0.4(default)      raxml/7.2.8
garli/0.96b8(default)     mrbayes/3.1.2-mpi          raxml/7.2.7
```

```
----- /bubo/sw/mf/kalkyl/bioinfo-tools/pipelines -----
ab_wtp/1.1(default)       cufflinks/0.9.2            cufflinks/1.1.0           tophat/1.2.0
bowtie/0.12.6(default)    cufflinks/0.9.3            cufflinks/1.2.1           tophat/1.3.3
```



# UPPMAX Commands

## uquota

```
[dahlo@biologin work]$ uquota
```

Your File Area	Usage (GB)	Quota Limit (GB)	Over Quota	Grace Time
-----	-----	-----	-----	-----
dahlo glob	196	2048		-
dahlo home	4	32		-
/proj/b2010015	229	256		
/proj/b2010015/nobackup	0	512		-
/proj/b2010033	132	6348		
/proj/b2010033/nobackup	27	512		-

# UPPMAX Commands

## projinfo

```
[dahlo@kalkyl4 work]$ projinfo
(Counting the number of core hours used since 2012-08-19/00:00:00 until now.)
```

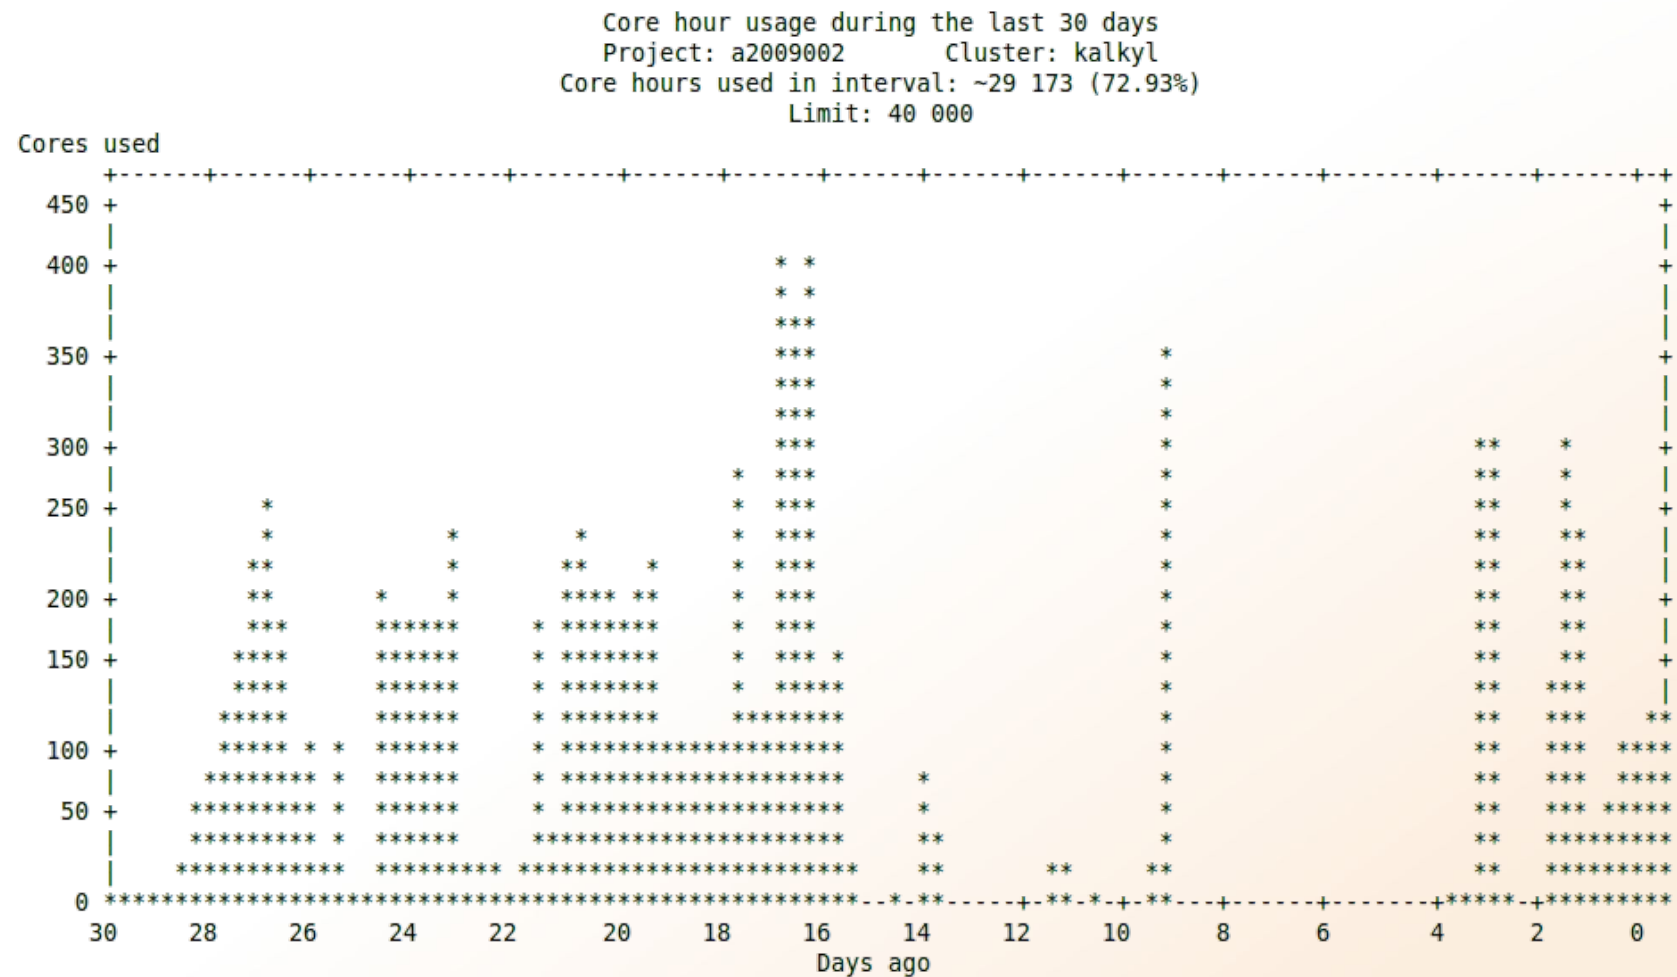
Project User	Used[h]	Current allocation [h/month]
-----		
b2010015	1257.20	2000
ameur	1257.20	
-----		
b2010069	0.00	2000
-----		
b2010074	110.98	2000
dahlo	1.01	
seba	109.97	
-----		
b2012044	0.00	2000
-----		
g2012005	0.00	2000
-----		
g2012083	0.00	2000
-----		
g2012157	0.12	2000
dahlo	0.12	

```
[dahlo@kalkyl4 work]$
```



# UPPMAX Commands

`projplot -A <proj-id> (-h for more options)`




# UPPMAX Commands


Plot efficiency

```
jobstats -p -A <projid>
```


Flags: mem\_underused:126:0



Flags: mem\_underused:504.7:7.9 node\_type\_misbooked:mem512GB:mem128GB



Flags: none



## Summary

All jobs are run on nodes through queue system

A job script usually consists of

- Job settings (-A, -p, -n, -t)

- Modules to be loaded

- Bash code to perform actions

  - Run a program, or multiple programs

## More info on UPPMAX homepage

<http://www.uppmax.uu.se/milou-user-guide>

Laboratory time! (again)